

Properties of MLE: Uncertainty in Parameters

INFO 521 · Module 3 · Lecture B — Properties of the MLE

Greg Chism

College of Information Science · University of Arizona

August 2026

Adapted with permission from lecture materials by Clayton Morrison, University of Arizona.

Learning outcomes

By the end of this module you should be able to:

1. **Show the MLE solution is unique** — that the log-likelihood has a single maximum (a negative-definite Hessian).
2. **State what unbiasedness of $\hat{\mathbf{w}}$ does and does not mean** — and the common claims it is *not*.
3. **Quantify parameter uncertainty** — via $\text{Cov}[\hat{\mathbf{w}}]$ and the Fisher information.

Recap: likelihood and the MLE

The linear-Gaussian likelihood:

$$p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}, \sigma^2) = \prod_{n=1}^N \mathcal{N}(y_n \mid \mathbf{w}^\top \mathbf{x}_n, \sigma^2)$$

Its maximizers — the MLE (derived in the likelihood lecture of this module):

$$\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}, \quad \hat{\sigma}^2 = \frac{1}{N} (\mathbf{y}^\top \mathbf{y} - \mathbf{y}^\top \mathbf{X} \hat{\mathbf{w}})$$

**Unique, unbiased, and
how uncertain.**

Property 1 · Uniqueness

The log-likelihood — we write it $\ell(\mathbf{w}) = \log p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}, \sigma^2)$, reserving $\mathcal{L}$ for losses:

$$\ell(\mathbf{w}) = -\frac{N}{2} \log 2\pi - N \log \sigma - \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \mathbf{w}^\top \mathbf{x}_n)^2$$

Gradient — vanishes at the normal equations:


$$\frac{\partial \ell}{\partial \mathbf{w}} = \frac{1}{\sigma^2} (\mathbf{X}^\top \mathbf{y} - \mathbf{X}^\top \mathbf{X} \mathbf{w})$$

Hessian — constant and **negative definite** ($\mathbf{v}^\top \mathbf{H} \mathbf{v} < 0 \ \forall \mathbf{v} \neq \mathbf{0}$):

$$\mathbf{H} = \frac{\partial^2 \ell}{\partial \mathbf{w} \partial \mathbf{w}^\top} = -\frac{1}{\sigma^2} \mathbf{X}^\top \mathbf{X} \implies \text{a single, unique maximum.}$$

The same $\mathbf{X}^\top \mathbf{X}$ full-column-rank condition from Module 1 — now reappearing as a negative-definite Hessian.

The generative picture



Expectation — a reminder

For a random vector $\mathbf{x}$ with density $p(\mathbf{x})$:

$$\mathbb{E}_{p(\mathbf{x})}[f(\mathbf{x})] = \int f(\mathbf{x}) p(\mathbf{x}) d\mathbf{x}$$

Mean $\boldsymbol{\mu} = \mathbb{E}[\mathbf{x}]$

Covariance $\text{Cov}[\mathbf{x}] = \mathbb{E}[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^\top]$

Property 2 · Unbiasedness

Take the expectation of $\hat{\mathbf{w}}$ over datasets from the generating process (only $\mathbf{y}$ is random; $\mathbf{X}$ is fixed):

$$\begin{aligned}\mathbb{E}[\hat{\mathbf{w}}] &= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbb{E}[\mathbf{y}] \\ &= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{X} \mathbf{w} \\ &= \mathbf{w} \quad \implies \quad \hat{\mathbf{w}} \text{ is unbiased.}\end{aligned}$$

What “unbiased” means — and doesn’t

Unbiased means: across repeated size- N samples from the generating process, the collection of estimates $\hat{\mathbf{w}}$ is **centered** on the true $\mathbf{w}$.

It is **not** the same as:

- “**more data $\Rightarrow$ better estimates**” — true here, but that is a *separate* sample-size effect, not what unbiasedness asserts;
- **UMVU** (uniformly minimum-variance unbiased) — a strictly stronger claim about being *closest*, which unbiasedness alone does not give;
- a guarantee of being best for *your* dataset — a **biased** estimator can land closer to $\mathbf{w}$ from *less* data.

Property 3 · Covariance

The punchline first — the parameter covariance is

$$\text{Cov}[\hat{\mathbf{w}}] = \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1}$$

Derived in full in the appendix → [appendix-covariance.html](#).

Fisher information

The (expected) Fisher information is the curvature of the log-likelihood:

$$\mathcal{I} = \mathbb{E} \left[- \frac{\partial^2 \log p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}, \sigma^2)}{\partial \mathbf{w} \partial \mathbf{w}^\top} \right] = \frac{1}{\sigma^2} \mathbf{X}^\top \mathbf{X} \quad \implies \quad \text{Cov}[\hat{\mathbf{w}}] = \mathcal{I}^{-1}$$

- **Diagonal** $\mathcal{I}_{jj}$ — information about each parameter on its own;
- **Off-diagonal** $\mathcal{I}_{jk}$ — information *shared* between a pair (their estimates trade off).

Parameter uncertainty on NHANES




Figure 1: $\text{Cov}[\hat{w}] = \sigma^2(X^T X)^{-1}$ shrinks as N grows: $N = 40$ (orange, dashed) vs. full N (green, solid).

Live demo: the parameter-uncertainty explorer

Draw datasets from a fixed process — watch $\hat{\mathbf{w}}$ scatter and fill the $\text{Cov}[\hat{\mathbf{w}}]$ ellipse, live on this slide.

Through-line

Uniqueness (negative-definite Hessian) → Fisher information $\mathcal{I} = \frac{1}{\sigma^2} \mathbf{X}^\top \mathbf{X}$ → (Module 5) the Laplace approximation's posterior covariance $\mathbf{H}^{-1} = \mathcal{I}^{-1}$.

The curvature of the log-likelihood is one object doing three jobs, three modules apart — **same** $\mathbf{X}^\top \mathbf{X}$, **same inverse**.

Recap + checkpoint preview

Three properties of the linear-Gaussian MLE:

1. **Unique** — $\mathbf{H} = -\frac{1}{\sigma^2} \mathbf{X}^\top \mathbf{X}$ is negative definite.
2. **Unbiased** — $\mathbb{E}[\hat{\mathbf{w}}] = \mathbf{w}$.
3. **Uncertain, quantifiably** — $\text{Cov}[\hat{\mathbf{w}}] = \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1} = \mathcal{I}^{-1}$.

Checkpoint quiz 1.2 (Week 3, D2L) gates milestone 1.2 on reproducing the **uniqueness** and **unbiasedness** arguments; the MLE == OLS check lives in Homework Unit 3.