

Properties of MLE: Uncertainty in Parameters

INFO 521 · Module 3

Greg Chism

School of Information · University of Arizona

June 2026

Learning outcomes

By the end of this module you should be able to:

1. **Show the MLE solution is unique** — that the log-likelihood has a single maximum (a negative-definite Hessian).
2. **State what unbiasedness of $\hat{\mathbf{w}}$ does and does not mean** — and the common claims it is *not*.
3. **Quantify parameter uncertainty** — via $\text{cov}\{\hat{\mathbf{w}}\}$ and the Fisher information.

Recap: likelihood and the MLE

The linear-Gaussian likelihood:

$$\mathcal{L} = p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}, \sigma^2) = \prod_{n=1}^N \mathcal{N}(y_n \mid \mathbf{w}^\top \mathbf{x}_n, \sigma^2)$$

Its maximizers — the MLE we derived last module:

$$\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}, \quad \hat{\sigma}^2 = \frac{1}{N} (\mathbf{y}^\top \mathbf{y} - \mathbf{y}^\top \mathbf{X} \hat{\mathbf{w}})$$

**Unique, unbiased, and
how uncertain.**

Property 1 • Uniqueness

The log-likelihood:

$$\log \mathcal{L} = -\frac{N}{2} \log 2\pi - N \log \sigma - \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \mathbf{w}^\top \mathbf{x}_n)^2$$

Gradient — vanishes at the normal equations:


$$\frac{\partial \log \mathcal{L}}{\partial \mathbf{w}} = \frac{1}{\sigma^2} (\mathbf{X}^\top \mathbf{y} - \mathbf{X}^\top \mathbf{X} \mathbf{w})$$

Hessian — constant and **negative definite** ($\mathbf{x}^\top \mathbf{H} \mathbf{x} < 0$):

$$\mathbf{H} = \frac{\partial^2 \log \mathcal{L}}{\partial \mathbf{w} \partial \mathbf{w}^\top} = -\frac{1}{\sigma^2} \mathbf{X}^\top \mathbf{X} \implies \text{a single, unique maximum.}$$

The same $\mathbf{X}^\top \mathbf{X}$ full-column-rank condition from Module 1 — now reappearing as a negative-definite Hessian.

The generative picture



Expectation — a reminder

For a random vector $\mathbf{x}$ with density $p(\mathbf{x})$:

$$\mathbb{E}_{p(\mathbf{x})}\{f(\mathbf{x})\} = \int f(\mathbf{x}) p(\mathbf{x}) d\mathbf{x}$$

Mean $\boldsymbol{\mu} = \mathbb{E}[\mathbf{x}]$

Covariance $\text{cov}\{\mathbf{x}\} = \mathbb{E}[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^\top]$

Property 2 · Unbiasedness

Take the expectation of $\hat{\mathbf{w}}$ over datasets from the generating process (only $\mathbf{y}$ is random; $\mathbf{X}$ is fixed):

$$\begin{aligned}\mathbb{E}[\hat{\mathbf{w}}] &= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbb{E}[\mathbf{y}] \\ &= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{X} \mathbf{w} \\ &= \mathbf{w} \quad \implies \quad \hat{\mathbf{w}} \text{ is unbiased.}\end{aligned}$$

What “unbiased” means — and doesn’t

Unbiased means: across repeated size- N samples from the generating process, the collection of estimates $\hat{\mathbf{w}}$ is **centered** on the true $\mathbf{w}$.

It is **not** the same as:

- “**more data $\Rightarrow$ better estimates**” — true here, but that is a *separate* sample-size effect, not what unbiasedness asserts;
- **UMVU** (uniformly minimum-variance unbiased) — a strictly stronger claim about being *closest*, which unbiasedness alone does not give;
- a guarantee of being best for *your* dataset — a **biased** estimator can land closer to $\mathbf{w}$ from *less* data.

Property 3 · Covariance

The punchline first — the parameter covariance is

$$\text{cov}\{\hat{\mathbf{w}}\} = \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1}$$

Derived in full in the appendix → [appendix-covariance.html](#).

Fisher information

The (expected) Fisher information is the curvature of the log-likelihood:

$$\mathcal{I} = \mathbb{E} \left[- \frac{\partial^2 \log p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}, \sigma^2)}{\partial \mathbf{w} \partial \mathbf{w}^\top} \right] = \frac{1}{\sigma^2} \mathbf{X}^\top \mathbf{X} \quad \implies \quad \text{cov}\{\hat{\mathbf{w}}\} = \mathcal{I}^{-1}$$

- **Diagonal** $\mathcal{I}_{jj}$ — information about each parameter on its own;
- **Off-diagonal** $\mathcal{I}_{jk}$ — information *shared* between a pair (their estimates trade off).

Parameter uncertainty on NHANES




Figure 1: $\text{cov}\{\hat{w}\} = \hat{\sigma}^2(X^T X)^{-1}$ shrinks as N grows: $N = 40$ (orange, dashed) vs. full N (green, solid).

Live demo: parameter-uncertainty explorer

Pipeline test — placeholder tool; the parameter-uncertainty viz is a future M3 build.

Example Domain

This domain is for use in documentation examples without needing permission.
Avoid use in operations.

[Learn more](#)

Explorer — static fallback

The iframe above does not print or export to PDF; this snapshot stands in.




Figure 2: Static fallback: the covariance ellipses (used in print / PDF).

Through-line

Uniqueness (negative-definite Hessian) $\rightarrow$ Fisher information $\mathcal{I} = \frac{1}{\sigma^2} \mathbf{X}^\top \mathbf{X} \rightarrow$ (Module 5) the Laplace approximation's posterior covariance $\mathbf{H}^{-1} = \mathcal{I}^{-1}$.

The curvature of the log-likelihood is one object doing three jobs, three modules apart — **same $\mathbf{X}^\top \mathbf{X}$, same inverse.**

Recap + checkpoint preview

Three properties of the linear-Gaussian MLE:

1. **Unique** — $\mathbf{H} = -\frac{1}{\sigma^2} \mathbf{X}^\top \mathbf{X}$ is negative definite.
2. **Unbiased** — $\mathbb{E}[\hat{\mathbf{w}}] = \mathbf{w}$.
3. **Uncertain, quantifiably** — $\text{cov}\{\hat{\mathbf{w}}\} = \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1} = \mathcal{I}^{-1}$.

The **P1-M2 closed-book checkpoint** gates on reproducing the **uniqueness** and **unbiasedness** arguments from scratch.